

1. Использовал docker-образы Bitnami для установки обеих систем: Kafka и Spark, т.к. установить непосредственно по инструкции, показанной на лекции, не получилось.

ПО	Версия ПО	Ссылка на Git	Команда docker
Kafka	3.8.0	https://github.com/bitnami/containers/tree/main/bitnami/kafka	docker pull bitnami/kafka
Spark	3.5.2	https://github.com/bitnami/containers/tree/main/bitnami/spark	docker pull bitnami/spark

2. Создал **docker-compose.yml** файл для запуска контейнеров в единой сети:

```
services:
  zookeeper:
    image: bitnami/zookeeper
    ports:
      - "2181:2181"
    environment:
      ALLOW_ANONYMOUS_LOGIN: "yes"
    networks:
      - kafka-network
  kafka:
    image: bitnami/kafka
    environment:
      KAFKA_CFG_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_LISTENERS:
        INSIDE://PLAINTEXT://0.0.0.0:29092,OUTSIDE://PLAINTEXT://0.0.0.0:9092
      KAFKA_ADVERTISED_LISTENERS: INSIDE://kafka:29092,OUTSIDE://172.21.182.15:9092
      KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
        INSIDE:PLAINTEXT,OUTSIDE:PLAINTEXT
      KAFKA_INTER_BROKER_LISTENER_NAME: INSIDE
    ports:
      - "9092:9092"
    depends_on:
      - zookeeper
    networks:
      - kafka-network
  spark-master:
    image: bitnami/spark
    ports:
      - "8080:8080"
      - "7077:7077"
    volumes:
      - /mnt/d/Курсы и тренинги/Data Engineer Нетология/Spark Structured Streaming:/opt/spark/work-dir
    environment:
      - SPARK_MODE=master
    networks:
      - kafka-network
  spark-worker:
    image: bitnami/spark
    volumes:
      - /mnt/d/Курсы и тренинги/Data Engineer Нетология/Spark Structured Streaming:/opt/spark/work-dir
```

```
environment:
- SPARK_MODE=worker
- SPARK_WORKER_MASTER=spark://spark-master:7077
depends_on:
- spark-master
networks:
- kafka-network
networks:
kafka-network:
driver: bridge
```

Здесь добавлена сеть **kafka-network** для того, чтобы оба контейнера видели друг друга и могли обмениваться данными.

Добавление Zookeeper было необязательным, просто «из коробки» Kafka запускалась с ошибками и, потратив много сил, нашел в сети решение, основанное на Zookeeper. Сейчас понимаю, что можно было запустить без Zookeeper, в режиме KRaft.

Что важно: пришлось переписать настройки Listeners, т.к. исходные Listeners от Bitnami настроены так, что не допускают подключения внешних клиентов:

- KAFKA_CFG_LISTENERS=PLAINTEXT://:9092,CONTROLLER://:9093
- KAFKA_CFG_ADVERTISED_LISTENERS=PLAINTEXT://:9092

В **docker-compose.yml** я разделил Listeners на внутренних и внешних, через указание разных портов, на которых Kafka слушает внутренних и внешних клиентов.

В качестве IP для внешнего хоста использовал IP WSL своего ПК ('172.21.182.15' определил командой `ip addr` в Bash)

Также в контейнеры добавлена опция **volumes** с указанием пути до python и spark-скриптов на моем локальном ПК для монтирования томов и обеспечения возможности запуска скриптов непосредственно внутри контейнеров.

3. Выполнил в терминале `docker-compose up`, развернул контейнеры в единой сети
4. В своем Python-окружении установил необходимые библиотеки:

- `pip install kafka`
- `pip install pyspark`

5. Поскольку Kafka установлена в контейнере, пришлось внести правки в скрипт **producer.py**:

- вместо строки `KafkaProducer(bootstrap_servers=['localhost:9092'])` прописал `KafkaProducer(bootstrap_servers=['172.21.182.15:9092'])`, где '172.21.182.15' – мой IP WSL (определил командой `ip addr` в Bash)
- в строке `producer.send` добавил методы `.add_callback` и `.add_errback`, вызывающие соответствующие функции для формирования логов и получения обратной связи о том, было ли сообщение успешно отправлено:

```
# Функция для обратного вызова при успешной отправке сообщения
def on_send_success(record_metadata):
    print(f'Message sent to {record_metadata.topic} partition
{record_metadata.partition} with offset {record_metadata.offset}')

# Функция для обратного вызова при ошибке
def on_send_error(excp):
```

```
print(f"Message delivery failed: {excp}")
```

6. Конечная версия скрипта **producer.py**, генерирующего входной поток для Kafka:

```
from time import sleep
from json import dumps
from kafka import KafkaProducer
from random import randrange, choices
import string

# Функция для обратного вызова при успешной отправке сообщения
def on_send_success(record_metadata):
    print(f"Message sent to {record_metadata.topic} partition {record_metadata.partition} with offset {record_metadata.offset}")

# Функция для обратного вызова при ошибке
def on_send_error(excp):
    print(f"Message delivery failed: {excp}")

producer = KafkaProducer(bootstrap_servers=['172.21.182.15:9092'],
                        value_serializer=lambda x: dumps(x).encode('utf-8'),
                        retries=5)

def push():
    for e in range(50):
        text = ".join(choices(string.ascii_uppercase + string.digits, k=20))
        user = {'id': randrange(5), 'action': text}
        # Отправляем сообщение с обработчиками
        producer.send('netology-spark',
value=user).add_callback(on_send_success).add_errback(on_send_error)
        sleep(randrange(3))
        print("Produced message ", e)

try:
    while True:
        push()
except KeyboardInterrupt:
    producer.close()
    print("Exit")
```

7. Запуск скрипта генерирует принты вида:

```
Message sent to netology-spark partition 0 with offset 255
Produced message 0
Produced message 1
Message sent to netology-spark partition 0 with offset 256
Message sent to netology-spark partition 0 with offset 257
Produced message 2
Message sent to netology-spark partition 0 with offset 258
Produced message 3
Message sent to netology-spark partition 0 with offset 259
Produced message 4
Produced message 5
Message sent to netology-spark partition 0 with offset 260
```

Message sent to netology-spark partition 0 with offset 261

которые свидетельствует об успешной отправке сообщений в Kafka.

8. Заходим в контейнер Kafka. Для этого:

- командой `docker ps` смотрим id контейнера Kafka
- вводим команду `docker exec -it <id контейнера> /bin/bash`
- для чтения входящего потока внутри контейнера вводим `kafka-console-consumer.sh -`
`-bootstrap-server kafka:9092 --topic netology-spark --from-beginning`
- видим поступающие от продюсера сообщения в формате бинарных данных:

```
quantum@DESKTOP-VJG26RT: ~
I have no name!@cac5527e8ad2:/$ kafka-console-consumer.sh --bootstrap-server kafka:9092 --topic netology-spark --from-beginning
^C
that's me
^C
hi
I can't believe - it's working?
^C
^C
^C
{"id": 0, "action": "UBUE3LV9W3MJCJNN03C9"}
{"id": 2, "action": "7YHDJG6IIWVBLPL0TSU5"}
{"id": 2, "action": "QFEI85SNVGH950DXQDQP"}
{"id": 4, "action": "0QSWHTEHLMDCXYCFX7AD"}
{"id": 1, "action": "RJBNEWD2A4FX3YGD98IC"}
{"id": 3, "action": "C3KZFQW952RSHY4WE94C"}
{"id": 1, "action": "A6WVWVHNLDC1VNT0FCWT"}
{"id": 4, "action": "084Y4JOP5GG7CRL8QVCV"}
{"id": 4, "action": "B8F4DKRSM4XLTC LMV1JT"}
{"id": 1, "action": "ZWAF0VKUDVIZMSV11P5F"}
{"id": 2, "action": "7J4BPZEHG7QE707B7MP"}
{"id": 3, "action": "QV3SH0CPGJ03H9LJC3G9"}
{"id": 4, "action": "514T0ZTJE9I0FEFI773H"}
{"id": 3, "action": "I4HK4T23RL0Q4Y278ESJ"}
{"id": 2, "action": "JQW9G1XNUTVQVCYXVFY7"}
{"id": 0, "action": "PLDH9KVKYN6XM37Y7JAT"}
{"id": 0, "action": "B91NXVBIQDT00Y5PUB7H"}
{"id": 4, "action": "NQIWI8KUYR1G84YUP9AD"}
{"id": 1, "action": "TH6P3STIUEUKNAHB1XAX"}
{"id": 2, "action": "SRB8J5V8QQFIHNKIYP8V"}
{"id": 2, "action": "LOWL33MJJ48PNHJ8ABPL"}
{"id": 2, "action": "50DEMXXNMAE87S91M114T"}
{"id": 1, "action": "HW5I0239EUSAGYTRP2H9"}
{"id": 3, "action": "1LHDC604AZGPGSXYBDBY"}
{"id": 4, "action": "HXPINZVMNB8FOWX609UY"}
{"id": 4, "action": "03IDL0N11MBSSAK7NU9W"}
{"id": 2, "action": "K8I6FQ8RKOEPW99HA30Y"}
{"id": 4, "action": "UORJBGDLEIJ15XRI46GZ"}
{"id": 4, "action": "SJOM85805UXICUCOEQR1"}
{"id": 2, "action": "T2JSY3DSYZ22ESMQQ7XA"}
{"id": 3, "action": "I6264TLETSNWZNRMPDGZ"}
{"id": 0, "action": "OFR1X4S612EIH3CM35P"}
{"id": 0, "action": "48X7H0IAHP3C6YKU9LHB"}
{"id": 1, "action": "YOWRL0PY74F320Q6DP6Y"}
{"id": 3, "action": "0RXVHM7J660CC2JB5RJ3"}
{"id": 4, "action": "R6Q0ZK65MTRYJA1JJ1ZI"}
{"id": 4, "action": "PGGK06C8I938ND323WOP"}
{"id": 4, "action": "XUDI5Z0GBSMGNMQM487Q"}
{"id": 1, "action": "LAYQ40KGCVCQ8S3FP1NFI"}
^C
```

9. Далее заходим в контейнер Spark с помощью команды

`docker exec -t <id контейнера Spark> /bin/bash`

10. Для получения входного потока из Kafka внутри контейнера Spark выполняем
`/opt/bitnami/spark/conf$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-`
`10_2.12:3.5.2 /opt/spark/work-dir/structure_streaming_kafka_test.py`

Я создал бэкап версию Spark-скрипта (`structure_streaming_kafka_test.py`) и работал с ней.
Данный скрипт хранится в директории, к которой был смонтирован том в разделе `volumes`
`docker-compose.yml`

11. Дальнейшие манипуляции и распаковка данных из Kafka осуществлялись исключительно
через редактирование файла `structure_streaming_kafka_test.py`:

- Чтение бинарных данных с логированием на уровне **info**:

```
quantum@DESKTOP-VJG26RT: ~
24/09/20 18:52:22 INFO SubscriptionState: [Consumer clientId=consumer-spark-kafka-source-9557a452-11fc-4565-bbf1-cb578858efc7--1060461563-executor-1, groupId=spark-kafka-source-9557a452-11fc-4565-bbf1-cb578858efc7--1060461563-executor-1] Seeking to latest offset of partition netology-spark-0
24/09/20 18:52:22 INFO SubscriptionState: [Consumer clientId=consumer-spark-kafka-source-9557a452-11fc-4565-bbf1-cb578858efc7--1060461563-executor-1, groupId=spark-kafka-source-9557a452-11fc-4565-bbf1-cb578858efc7--1060461563-executor-1] Resetting offset for partition netology-spark-0 to position FetchPosition{offset=68, offsetEpoch=Optional.empty, currentLeader=LeaderAndEpoch{leader=Optional[172.21.182.100], epoch=0}}, epoch=0}}.
24/09/20 18:52:22 INFO DataWritingSparkTask: Writer for partition 0 is committing.
24/09/20 18:52:22 INFO DataWritingSparkTask: Committed partition 0 (task 3, attempt 0, stage 3.0)
24/09/20 18:52:22 INFO KafkaDataConsumer: From Kafka topicPartition=netology-spark-0 groupId=spark-kafka-source-9557a452-11fc-4565-bbf1-cb578858efc7--1060461563-executor-1 read 1 records through 1 records), taking 508572247 nanos, during time span of 509415987 nanos.
24/09/20 18:52:22 INFO Executor: Finished task 0.0 in stage 3.0 (TID 3). 2242 bytes result sent to driver
24/09/20 18:52:22 INFO TaskSetManager: Finished task 0.0 in stage 3.0 (TID 3) in 529 ms on 410e78e0046b (executor driver) (1/1)
24/09/20 18:52:22 INFO TaskSchedulerImpl: Removed TaskSet 3.0, whose tasks have all completed, from pool
24/09/20 18:52:22 INFO DAGScheduler: ResultStage 3 (start at NativeMethodAccessorImpl.java:0) finished in 0.539 s
24/09/20 18:52:22 INFO DAGScheduler: Job 3 is finished. Cancelling potential speculative or zombie tasks for this job
24/09/20 18:52:22 INFO TaskSchedulerImpl: Killing all running tasks in stage 3: Stage finished
24/09/20 18:52:22 INFO DAGScheduler: Job 3 finished: start at NativeMethodAccessorImpl.java:0, took 0.542072 s
24/09/20 18:52:22 INFO WriteToDataSourceV2Exec: Data source write support MicroBatchWrite[epoch: 3, writer: ConsoleWriter[numRows=20, truncate=true]] is committing.
-----
Batch: 3
-----
+-----+-----+-----+-----+-----+-----+
| key|          value|      topic|partition|offset|      timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+
|NULL|[7B 22 69 64 22 3...]|netology-spark|      0|    67|2024-09-20 18:52:...|      0|
+-----+-----+-----+-----+-----+-----+

24/09/20 18:52:22 INFO WriteToDataSourceV2Exec: Data source write support MicroBatchWrite[epoch: 3, writer: ConsoleWriter[numRows=20, truncate=true]] committed.
24/09/20 18:52:22 INFO CheckpointFileManager: Writing atomically to file:/tmp/temporary-1f0c8e51-1f8c-4815-84eb-83a769a10c87/commits/3 using temp file file:/tmp/temporary-1f0c8e51-1f8c-4815-84eb-mits/.3.52ca9b11-e5ee-4856-855e-2e25cabe7564.tmp
24/09/20 18:52:22 INFO CheckpointFileManager: Renamed temp file file:/tmp/temporary-1f0c8e51-1f8c-4815-84eb-83a769a10c87/commits/.3.52ca9b11-e5ee-4856-855e-2e25cabe7564.tmp to file:/tmp/temporary-815-84eb-83a769a10c87/commits/3
24/09/20 18:52:22 INFO MicroBatchExecution: Streaming query made progress: {
  "id" : "6766a304-0345-48d7-9708-fe9fe345ff97",
  "runId" : "8620aac1-8636-4a61-9af9-0e8f7e2005ed",
  "name" : null,
  "timestamp" : "2024-09-20T18:52:22.143Z",
  "batchId" : 3,
  "numInputRows" : 1,
  "inputRowsPerSecond" : 71.42857142857143,
  "processedRowsPerSecond" : 1.497005988023952,
  "durationMs" : {
    "addBatch" : 598,
    "commitOffsets" : 27,
    "getBatch" : 0,
    "latestOffset" : 4,
    "queryPlanning" : 8,
    "triggerExecution" : 668,
    "walCommit" : 29
  },
  "stateOperators" : [ ],
  "sources" : [ {
    "description" : "KafkaV2[Subscribe[netology-spark]]",
    "startOffset" : {
      "netology-spark" : {
```

- Чтение бинарных данных с логированием на уровне **error**:

```
quantum@DESKTOP-VJG26RT: ~
0 artifacts copied, 11 already retrieved (0kB/9ms)
-----
Batch: 0
-----
+-----+-----+-----+-----+-----+-----+-----+
|key|value|topic|partition|offset|timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+-----+
-----

Batch: 1
-----
+-----+-----+-----+-----+-----+-----+-----+
| key|          value|          topic|partition|offset|          timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+-----+
|NULL|[7B 22 69 64 22 3...|netology-spark|      0|  116|2024-09-20 19:19:...|          0|
+-----+-----+-----+-----+-----+-----+-----+
-----

Batch: 2
-----
+-----+-----+-----+-----+-----+-----+-----+
| key|          value|          topic|partition|offset|          timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+-----+
|NULL|[7B 22 69 64 22 3...|netology-spark|      0|  117|2024-09-20 19:19:...|          0|
+-----+-----+-----+-----+-----+-----+-----+
-----

Batch: 3
-----
+-----+-----+-----+-----+-----+-----+-----+
| key|          value|          topic|partition|offset|          timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+-----+
|NULL|[7B 22 69 64 22 3...|netology-spark|      0|  118|2024-09-20 19:19:...|          0|
+-----+-----+-----+-----+-----+-----+-----+
-----

Batch: 4
-----
+-----+-----+-----+-----+-----+-----+-----+
| key|          value|          topic|partition|offset|          timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+-----+
|NULL|[7B 22 69 64 22 3...|netology-spark|      0|  119|2024-09-20 19:19:...|          0|
+-----+-----+-----+-----+-----+-----+-----+
-----

Batch: 5
-----
+-----+-----+-----+-----+-----+-----+-----+
| key|          value|          topic|partition|offset|          timestamp|timestampType|
+-----+-----+-----+-----+-----+-----+-----+
|NULL|[7B 22 69 64 22 3...|netology-spark|      0|  120|2024-09-20 19:19:...|          0|
+-----+-----+-----+-----+-----+-----+-----+
```

- Распакованные данные после применения схемы и перевода в json:

quantum@DESKTOP-VJG26RT: ~

Batch: 0

```

+-----+
|timestamp|parsed_value|
+-----+
+-----+

```

Batch: 1

```

+-----+
|timestamp          |parsed_value          |
+-----+
|2024-09-20 19:30:24.815|{2, YU9L63IV4E8GSQNSPA80}|
|2024-09-20 19:30:24.816|{3, HG3BV4ZRI5Z9JJ7IX936}|
|2024-09-20 19:30:24.816|{2, J16NFBZWRL4NGG4KP37}|
+-----+

```

Batch: 2

```

+-----+
|timestamp          |parsed_value          |
+-----+
|2024-09-20 19:30:26.816|{1, VXU6RC48SZ94B05PAIEA}|
+-----+

```

Batch: 3

```

+-----+
|timestamp          |parsed_value          |
+-----+
|2024-09-20 19:30:27.817|{3, Z6MF0ZHNQL8553DGBEUN}|
+-----+

```

Batch: 4

```

+-----+
|timestamp          |parsed_value          |
+-----+
|2024-09-20 19:30:29.818|{2, Q056HQI4RAS1B82FLX5R}|
|2024-09-20 19:30:29.819|{4, P6370W4KEKSRVRHIK10M}|
+-----+

```

Batch: 5

```

+-----+
|timestamp          |parsed_value          |
+-----+
|2024-09-20 19:30:31.821|{4, 3KTIU2S5PC9GMUS1FW94}|
|2024-09-20 19:30:31.821|{4, 76M00P6DG756MV000Y601}|
+-----+

```

- Распакованные данные до уровня user_id и action:

quantum@DESKTOP-VJG26RT: ~

Batch: 0

timestamp	id	action
-----------	----	--------

Batch: 1

timestamp	id	action
2024-09-20 19:36:14.804	2	D9JZ94BFMSQLKA66ZXEU

Batch: 2

timestamp	id	action
2024-09-20 19:36:16.806	1	JM6CA1WC228TFIMHC441

Batch: 3

timestamp	id	action
2024-09-20 19:36:18.807	4	43AP3M8FUJT900KYFFZM

Batch: 4

timestamp	id	action
2024-09-20 19:36:19.808	4	YZFCL7M345SFG6RPY9LI

Batch: 5

timestamp	id	action
2024-09-20 19:36:20.809	4	2UC8VH9WNZ3DT4VHF3DL

- Данные после джойна входящего и статического потоков:

quantum@DESKTOP-VJG26RT: ~

Batch: 1

user_name	user_age	action	timestamp
Jimmy	18	GX626D0432906W285X15	2024-09-20 20:58:35.393
Jimmy	18	ADZ3IH8MPXSQKZIE653Q	2024-09-20 20:58:40.398
Jimmy	18	RHYH57UNVE7IFT8CBD4V	2024-09-20 20:58:40.398
Jimmy	18	BHYPW4PI8NPNWCETJN17	2024-09-20 20:59:06.504
Johnny	9	TXT6AG71QIZAFPPT3CR6	2024-09-20 20:59:06.503
Erle	40	YF76015GD4ACAYVNR02X	2024-09-20 20:59:06.503
Hank	48	CB9B17XZQS3797L23FUZ	2024-09-20 20:58:32.391
Hank	48	TNDNXEOHAHKSK4K3L744	2024-09-20 20:58:33.392
Hank	48	2V6HBGK1SASMJ8MDAUVN	2024-09-20 20:58:39.397
NULL	NULL	SG9RPPHJREPZDO3UWATR	2024-09-20 20:58:35.394
NULL	NULL	QRUMZWCUUM4KFFMYE5I7	2024-09-20 20:58:36.394
NULL	NULL	B3375QISWKINVB41XGA8	2024-09-20 20:58:38.396
NULL	NULL	S0HJ2HRQP7H8IEFY5JCY	2024-09-20 20:58:38.396
NULL	NULL	INR9Y3BIJSK1FCYJ1QD1	2024-09-20 20:58:39.397
NULL	NULL	VDITVQW4CMLMNSQUGWL1	2024-09-20 20:58:42.4
NULL	NULL	409YGXVCAO43FU6UA79H	2024-09-20 20:58:43.402

Batch: 2

user_name	user_age	action	timestamp
Johnny	9	3R7E5BB6X318GJTOIF2H	2024-09-20 20:59:08.505
Johnny	9	4QCH9XKCKV8BAQPLGZE0	2024-09-20 20:59:09.506
Johnny	9	AYST60KG1LTK5BZHT87I	2024-09-20 20:59:12.509
Johnny	9	790WKCS2IOX72Q2V2BNO	2024-09-20 20:59:12.51
Erle	40	5IUE2ZF54BFB5X5SRZE4	2024-09-20 20:59:11.508
Erle	40	OJU06689WFC8AY7GI9EX	2024-09-20 20:59:12.51
Hank	48	P35AA0CBI1XK6P980UHA	2024-09-20 20:59:08.505
Hank	48	J6L50IK4X6A8B4NAW3TP	2024-09-20 20:59:12.509

Batch: 3

user_name	user_age	action	timestamp
Johnny	9	WB2VXL42NEJQQHZWG3NY	2024-09-20 20:59:13.514
Johnny	9	PDHXT74RBLBTMKE6UIT	2024-09-20 20:59:14.517
Erle	40	IAXMWRDVIYVBXKSQ0W7N	2024-09-20 20:59:14.519
NULL	NULL	8KW89NGLVAHRZ5W58H3H	2024-09-20 20:59:13.51

- После добавления чекпойнтов фиксируем, что чтение данных из входного потока возобновляется не с 0, а с последнего прочитанного батча:

```
quantum@DESKTOP-VJG26RT: ~
I have no name!@410e78e0046b:/opt/bitnami/spark/conf$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.2 /opt/spark/work-dir/structure_streaming_kafka_test.py
:: loading settings :: url = jar:file:/opt/bitnami/spark/jars/ivy-2.5.1.jar!/org/apache/ivy/core/settings/ivysettings.xml
Ivy Default Cache set to: /opt/bitnami/spark/.ivy2/cache
The jars for the packages stored in: /opt/bitnami/spark/.ivy2/jars
org.apache.spark#spark-sql-kafka-0-10_2.12 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-5105eb00-7dab-49a2-b277-1cbb2223a2d1;1.0
  confs: [default]
  found org.apache.spark#spark-sql-kafka-0-10_2.12;3.5.2 in central
  found org.apache.spark#spark-token-provider-kafka-0-10_2.12;3.5.2 in central
  found org.apache.kafka#kafka-clients;3.4.1 in central
  found org.lz4#lz4-java;1.8.0 in central
  found org.xerial.snappy#snappy-java;1.1.10.5 in central
  found org.slf4j#slf4j-api;2.0.7 in central
  found org.apache.hadoop#hadoop-client-runtime;3.3.4 in central
  found org.apache.hadoop#hadoop-client-api;3.3.4 in central
  found commons-logging#commons-logging;1.1.3 in central
  found com.google.code.findbugs#jsr305;3.0.0 in central
  found org.apache.commons#commons-pool2;2.11.1 in central
:: resolution report :: resolve 370ms :: artifacts dl 12ms
  :: modules in use:
    com.google.code.findbugs#jsr305;3.0.0 from central in [default]
    commons-logging#commons-logging;1.1.3 from central in [default]
    org.apache.commons#commons-pool2;2.11.1 from central in [default]
    org.apache.hadoop#hadoop-client-api;3.3.4 from central in [default]
    org.apache.hadoop#hadoop-client-runtime;3.3.4 from central in [default]
    org.apache.kafka#kafka-clients;3.4.1 from central in [default]
    org.apache.spark#spark-sql-kafka-0-10_2.12;3.5.2 from central in [default]
    org.apache.spark#spark-token-provider-kafka-0-10_2.12;3.5.2 from central in [default]
    org.lz4#lz4-java;1.8.0 from central in [default]
    org.slf4j#slf4j-api;2.0.7 from central in [default]
    org.xerial.snappy#snappy-java;1.1.10.5 from central in [default]
  -----
  |               |          modules          |      artifacts      |
  |   conf   | number | search | dwnlded | evicted || number | dwnlded |
  |-----|-----|-----|-----|-----|
  | default |    11  |    0   |    0    |    0    ||    11  |    0    |
  -----
:: retrieving :: org.apache.spark#spark-submit-parent-5105eb00-7dab-49a2-b277-1cbb2223a2d1
  confs: [default]
  0 artifacts copied, 11 already retrieved (0kB/9ms)

Batch: 4
-----
+-----+-----+-----+-----+
|user_name|user_age|action          |timestamp          |
+-----+-----+-----+-----+
|Johnny   |9       |8TQXBQ4AHFYEWHA|2024-09-20 20:59:15.52|
+-----+-----+-----+-----+

Batch: 5
-----
+-----+-----+-----+-----+
|user_name|user_age|action          |timestamp          |
+-----+-----+-----+-----+
|Jimmy    |18      |TGQXMRXSAJRVCNG7|2024-09-20 20:59:17.522|
+-----+-----+-----+-----+
```

- Чтение данных после агрегации:

quantum@DESKTOP-VJG26RT: ~

Batch: 0

user_name	user_age	count
-----------	----------	-------

Batch: 1

user_name	user_age	count
Hank	48	1

Batch: 2

user_name	user_age	count
Jimmy	18	1
Johnny	9	1
Hank	48	1
NULL	NULL	1

Batch: 3

user_name	user_age	count
Jimmy	18	1
Johnny	9	1
Erle	40	1
Hank	48	2
NULL	NULL	2

Batch: 4

user_name	user_age	count
Jimmy	18	1
Johnny	9	2
Erle	40	1
Hank	48	2
NULL	NULL	3

12. Итоговый Spark-скрипт:

```
from pyspark.sql import SparkSession
from time import sleep
from pyspark.sql.functions import col, from_json
from pyspark.sql.types import StructType, StringType, IntegerType

# явным образом задаем структуру json-контента
schema = StructType().add("id", IntegerType()).add("action", StringType())

users_schema = StructType().add("id", IntegerType()).add("user_name",
StringType()).add("user_age", IntegerType())

spark = SparkSession.builder.appName("SparkStreamingKafka").getOrCreate()

input_stream = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "kafka:29092") \
    .option("subscribe", "netology-spark") \
    .option("failOnDataLoss", False) \
    .load()

#покажем входящий контент
#input_stream.writeStream.format("console").outputMode("append").start().awaitTermination()
#input_stream = input_stream.writeStream.format("console").outputMode("append").start()

json_stream = input_stream.select(col("timestamp").cast("string"),
from_json(col("value").cast("string"), schema).alias("parsed_value"))
#json_stream.writeStream.format("console").outputMode("append").option("truncate",
False).start().awaitTermination()

#выделим интересующие элементы
clean_data = json_stream.select(col("timestamp"), col("parsed_value.id").alias("id"),
col("parsed_value.action").alias("action"))
#clean_data.writeStream.format("console").outputMode("append").option("truncate",
False).start().awaitTermination()

#добавим join со статическим dataset - создаем данные
users_data = [(1, "Jimmy", 18), (2, "Hank", 48), (3, "Johnny", 9), (4, "Erle", 40)]
users = spark.createDataFrame(data=users_data, schema=users_schema)

#делаем join
#join_stream = clean_data.join(users, clean_data.id == users.id,
"left_outer").select(users.user_name, users.user_age, clean_data.action, clean_data.timestamp)
join_stream = clean_data.join(users, clean_data.id == users.id,
"left_outer").select(users.user_name, users.user_age, clean_data.action, clean_data.timestamp)
#join_stream.writeStream.format("console").outputMode("append").option("truncate",
False).start().awaitTermination()

#убираем terminate
#res_checkpoints= join_stream.writeStream.\
#format("console").\
#outputMode("append").\
```

```

#option("truncate", False).\
#option("checkpointLocation", "checkpoint/target").\
#start()

#sleep(10)
#res_checkpoints.stop()

#добавим агрегат - отображать число уникальных айдюков
stat_stream = clean_data.groupBy("id").count()
#stat_stream.writeStream.format("console").outputMode("complete").option("truncate",
False).start().awaitTermination()

join_stream_agg = stat_stream.join(users, stat_stream.id == users.id,
"left_outer").select(users.user_name, users.user_age, col('count'))

res_checkpoints_agg= join_stream_agg.writeStream.\
format("console").\
outputMode("complete").\
option("truncate", False).\
option("checkpointLocation", "checkpoint/target").\
start()\

sleep(10)
res_checkpoints_agg.stop()

```

13. Для сокращения количества логов внесена правка в конфигурационный файл **log4j2.properties** в контейнере Spark:

- Оригинальный файл **/opt/bitnami/spark/conf/log4j2.properties.template** скопирован; в копии убрано окончание **.template**
- В новом файле найден параметр **rootLogger.level = info**; **'info'** заменен на **'error'**

The screenshot shows the Docker Desktop interface for a container named 'quantum-spark-master-1'. The 'Files' tab is active, displaying a list of files in the container's filesystem. The file 'log4j2.properties' is highlighted, and its contents are shown in the editor below. The file is located at '/opt/bitnami/spark/conf/log4j2.properties'.

Name	Note	Size	Last modified	Mode
> bin			5 days ago	drwxr-xr-x
> conf	MODIFIED		4 hours ago	drwxrwxr-x
> checkpoint	ADDED		2 hours ago	drwxr-xr-x
fairscheduler.xml.template		1.1 kB	28 days ago	-rw-rw-r--
log4j2.properties	ADDED	3.3 kB	4 hours ago	-rw-rw-r--
log4j2.properties.template		3.3 kB	28 days ago	-rw-rw-r--
metrics.properties.template		8.9 kB	28 days ago	-rw-rw-r--
spark-defaults.conf	ADDED	1.3 kB	28 days ago	-rw-rw-r--
spark-env.sh	ADDED	53 B	7 hours ago	-rw-rw-r--

```

11 # Unless required by applicable law or agreed to in writing, software
12 # distributed under the License is distributed on an "AS IS" BASIS,
13 # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
14 # See the License for the specific language governing permissions and
15 # limitations under the License.
16 #
17
18 # Set everything to be logged to the console
19 rootLogger.level = error
20 rootLogger.appenderRef.stdout.ref = console
21

```